

# Housing Market Model (Team TNA)

```
In [16]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files in the input directory
import os

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
# You can also write temporary files to /kaggle/temp/, but they won't be saved
/kaggle/input/effects-of-policy-on-the-housing-market/sample_submissions.csv
/kaggle/input/effects-of-policy-on-the-housing-market/train.csv
/kaggle/input/effects-of-policy-on-the-housing-market/test.csv
```

## Uploading and splitting of data

```
In [17]: # Paths
input_path = '/kaggle/input/effects-of-policy-on-the-housing-market/'

# Load data # IDs removed from data # Names removed due to being relatively
train_df = pd.read_csv(os.path.join(input_path, 'train.csv'))
test_df = pd.read_csv(os.path.join(input_path, 'test.csv'))

X_train, y_train = train_df.drop(columns = train_df.columns[0:2]).drop(columns = train_df.columns[0:2])
```

## Exploratory Data Analysis

- `name` contains multiple values that should be split into a list to be placed into certain columns, and the first value appears to be categorical.

```
In [18]: X_train.head()
# float(re.sub("[^0-9.]", "", n))
```

Out [18]:

	id	name	neighborhood_overview	host_id	host_name
0	19792418	Home in Vancouver · ★4.75 · 1 bedroom · 1 bed ...	Everything you need is nearby.   Hig...	57488206	Jessi
1	1015650685503221866	Guest suite in Vancouver · ★New · 2 bedrooms · ...	NaN	139792573	Daniel
2	35265562	Guest suite in Vancouver · ★4.85 · 2 bedrooms ...	Beautiful neighbourhood close to prosperous Ma...	265504225	Alex
3	911948980885194155	Home in Vancouver · ★5.0 · 1 bedroom · 1 bed ...	We are located in a quiet residential neighbor...	22595056	Raymond
4	46069251	Guest suite in Vancouver · ★4.93 · 1 bedroom · ...	Kitsilano at it's best! Short walk to all the ...	65683877	Yendi

5 rows × 41 columns

## Data type for each column

- A `FunctionTransformer()` will be required to convert the data type for the categorical variables to string. The numerical variables should be scaled with a `StandardScaler()`.
- Some numerical variables are presented as strings rather than numbers (e.g. `host_response_rate`, `host_acceptance_rate`, `price`) and thus will need to be converted from their original data type with a different `FunctionTransformer()`.

In [19]: `X_train.dtypes`

```

Out[19]: id                int64
         name              object
         neighborhood_overview object
         host_id           int64
         host_name         object
         host_response_time object
         host_response_rate object
         host_acceptance_rate object
         host_is_superhost object
         host_listings_count int64
         host_total_listings_count int64
         neighbourhood      object
         neighbourhood_cleansed object
         latitude          float64
         longitude         float64
         property_type      object
         room_type          object
         accommodates       int64
         beds              float64
         amenities          object
         price              object
         minimum_nights     int64
         maximum_nights     int64
         minimum_nights_avg_ntm float64
         maximum_nights_avg_ntm float64
         availability_30     int64
         availability_60     int64
         availability_90     int64
         availability_365    int64
         number_of_reviews  int64
         number_of_reviews_ltm int64
         review_scores_rating float64
         review_scores_accuracy float64
         review_scores_cleanliness float64
         review_scores_checkin float64
         review_scores_communication float64
         review_scores_location float64
         review_scores_value float64
         instant_bookable   object
         calculated_host_listings_count int64
         reviews_per_month float64
dtype: object

```

## Testing the need for imputation within data

- Some of the data is missing, and thus imputation is required.
- `name` contains information on `bathrooms`, `bedrooms`, `beds` and `review_scores_rating`, while the `bathrooms` and `bedrooms` columns consist entirely of `NaN` data, so `name` will need to be modified to extract such information and allocate them accordingly.

```
In [20]: X_train.isna().any()
```

```

Out[20]: id                False
         name              False
         neighborhood_overview  True
         host_id           False
         host_name         False
         host_response_time  True
         host_response_rate  True
         host_acceptance_rate True
         host_is_superhost   True
         host_listings_count False
         host_total_listings_count False
         neighbourhood       True
         neighbourhood_cleansed False
         latitude            False
         longitude           False
         property_type       False
         room_type           False
         accommodates        False
         beds                True
         amenities           False
         price               True
         minimum_nights      False
         maximum_nights      False
         minimum_nights_avg_ntm False
         maximum_nights_avg_ntm False
         availability_30      False
         availability_60      False
         availability_90      False
         availability_365     False
         number_of_reviews    False
         number_of_reviews_ltm False
         review_scores_rating  True
         review_scores_accuracy True
         review_scores_cleanliness True
         review_scores_checkin True
         review_scores_communication True
         review_scores_location True
         review_scores_value  True
         instant_bookable     False
         calculated_host_listings_count False
         reviews_per_month    True
dtype: bool

```

## Categories within most categorical variables

- `name` is meant to be a categorical variable, but has not been transformed to allow for that until later.
- `host_response_time` and `host_is_superhost` will require imputation due to containing `nan`.
- `amenities` would theoretically require extraction of consistent features from `neighborhood_overview`.

```
In [21]: # print(X_train["name"].unique())
print(X_train["host_response_time"].unique()) # Missing values
print(X_train["host_is_superhost"].unique()) # Boolean, missing values
print(X_train["neighbourhood"].unique())
print(X_train["neighbourhood_cleansed"].unique())
print(X_train["property_type"].unique())
print(X_train["room_type"].unique())
# print(X_train["amenities"].unique()) # All empty
print(X_train["instant_bookable"].unique()) # Boolean

[nan 'within a few hours' 'within an hour' 'within a day'
 'a few days or more']
['f' 't' nan]
['Vancouver, British Columbia, Canada' nan
 'Delta, British Columbia, Canada'
 'Vancouver bc, British Columbia, Canada'
 'Vancouver , British Columbia, Canada' 'Vancouver, Canada'
 'West Vancouver, BC , Canada' 'Vancouver, British Columbia (BC), Canada']
['Hastings-Sunrise' 'Sunset' 'Riley Park' 'Kitsilano' 'Downtown'
 'Fairview' 'Dunbar Southlands' 'Kensington-Cedar Cottage' 'Marpole'
 'Mount Pleasant' 'West End' 'Grandview-Woodland' 'Renfrew-Collingwood'
 'Downtown Eastside' 'Oakridge' 'South Cambie' 'West Point Grey'
 'Arbutus Ridge' 'Killarney' 'Victoria-Fraserview' 'Kerrisdale'
 'Strathcona' 'Shaughnessy']
['Entire home' 'Entire guest suite' 'Entire condo' 'Entire rental unit'
 'Entire townhouse' 'Private room in guest suite' 'Private room in home'
 'Private room in townhouse' 'Private room in condo'
 'Private room in rental unit' 'Entire loft' 'Entire serviced apartment'
 'Shared room in rental unit' 'Private room in guesthouse'
 'Private room in villa' 'Room in boutique hotel' 'Entire guesthouse'
 'Casa particular' 'Shared room in home' 'Private room in tiny home'
 'Private room in loft' 'Tiny home' 'Entire timeshare'
 'Entire vacation home' 'Camper/RV' 'Tower' 'Entire cabin'
 'Room in aparthotel' 'Room in bed and breakfast'
 'Private room in casa particular' 'Private room in bed and breakfast'
 'Entire villa' 'Entire bungalow' 'Room in hotel' 'Entire place'
 'Shared room in condo' 'Entire cottage' 'Private room in bungalow'
 'Private room in boat' 'Private room in camper/rv' 'Boat'
 'Shared room in villa' 'Private room in serviced apartment'
 'Private room in castle' 'Cave' 'Shared room in loft']
['Entire home/apt' 'Private room' 'Shared room' 'Hotel room']
['f' 't']
```

## Most continuous variables

- From the data that can be seen, `beds` and several of the review score variables are missing significant amounts of data.
- Due to differing scales for most columns, `StandardScaler()` will be required.

```
In [22]: display(X_train.describe())
```

	id	host_id	host_listings_count	host_total_listings_count	
<b>count</b>	5.352000e+03	5.352000e+03	5352.000000	5352.000000	535
<b>mean</b>	4.398265e+17	1.916746e+08	11.917601	18.213191	4
<b>std</b>	4.317929e+17	1.808806e+08	45.576743	61.783016	
<b>min</b>	1.318800e+04	4.662000e+03	1.000000	1.000000	4
<b>25%</b>	3.365765e+07	2.637491e+07	1.000000	1.000000	4
<b>50%</b>	5.987246e+17	1.318506e+08	2.000000	3.000000	4
<b>75%</b>	8.769719e+17	3.428838e+08	5.000000	8.000000	4
<b>max</b>	1.044548e+18	5.505990e+08	466.000000	1037.000000	4

8 rows x 27 columns

## Extraction of Data and Preprocessing

### Methods to convert data within function transformers

```
In [23]: # Convert categorical column data type to string
from sklearn.preprocessing import FunctionTransformer
# https://stackoverflow.com/questions/59476179/is-it-possible-to-change-panda
def to_categorical(x):
    return pd.DataFrame(x).astype("string")
fun_tr = FunctionTransformer(to_categorical)
```

```
In [24]: from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler, FunctionTransformer
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.compose import make_column_transformer
from sklearn.feature_extraction.text import CountVectorizer

# Table transformation functions
def nameterms(x):
    if isinstance(x, (list)) or "." not in x:
        return x
    else:
        return x.split(".")

def strtoint(x, pos):
    if isinstance(x, list) and len(x) > abs(pos):
        new = x[pos]
    else:
        return np.nan
    n = re.sub("[^0-9.]", "", new)
    if n == "":
        return np.nan
    else:
```

```

        return float(n)

def strttoFloat(x):
    if isinstance(x, list) and x and len(x) != 4:
        rate = x[1]
    else:
        return np.nan
    ratenew = re.sub("[^0-9.]", "", str(rate))
    if ratenew == "":
        return np.nan
    else:
        return float(ratenew)

def strtopercent(x):
    return pd.DataFrame(x).map(lambda val: float(val.replace('%', '')) / 100)

def strtomoney(x):
    return pd.DataFrame(x).map(lambda val: float(val.replace('$', '').replace(

```

## Creation of `NameTransformer`

- Extract data from `name` column for use later in the table.
- `X_train_transformed` and `test_df_transformed` created to test transformers `NameTransformer` and later `preprocessor`.

```

In [25]: from sklearn.base import BaseEstimator, TransformerMixin

class NameTransformer(BaseEstimator, TransformerMixin):
    def fit(self, X, y=None):
        return self

    def transform(self, X):
        transformed_df = X.copy()
        transformed_df['name'] = transformed_df['name'].map(nameterms)
        transformed_df['bathrooms'] = transformed_df['name'].map(lambda x: s
        transformed_df['bedrooms'] = transformed_df['name'].map(lambda x: st
        transformed_df['beds'] = transformed_df['name'].map(lambda x: strtol
        transformed_df['review_scores_rating'] = transformed_df['name'].map(
        transformed_df['name'] = transformed_df['name'].apply(lambda x: x[0]
        return transformed_df

# Create the pipeline
pipeline_transformer = make_pipeline(
    NameTransformer(),
)

X_train_transformed = pipeline_transformer.fit_transform(X = X_train)
test_df_transformed = pipeline_transformer.fit_transform(test_df)

X_train_transformed.head()

```

Out [25]:

	id	name	neighborhood_overview	host_id	host_name
0	19792418	Home in Vancouver	Everything you need is nearby.   Hig...	57488206	Jessi
1	1015650685503221866	Guest suite in Vancouver	NaN	139792573	Daniel
2	35265562	Guest suite in Vancouver	Beautiful neighbourhood close to prosperous Ma...	265504225	Alex
3	911948980885194155	Home in Vancouver	We are located in a quiet residential neighbor...	22595056	Raymond
4	46069251	Guest suite in Vancouver	Kitsilano at it's best! Short walk to all the ...	65683877	Yendi

5 rows × 43 columns

## Preprocessing of features by column

- `percentage_features` and `currency_features` need to be converted from non-numeric values to numeric values. Due to `percentage_features` and `currency_features` being strings before transformation, their imputers both use "most\_frequent" as their imputation strategy.
- Since we did not know of any way to extract data from `neighborhood_overview` into `amenities`, both were dropped.
- `id`, `neighborhood_overview`, `host_id` and `host_name` are used more for identification than statistics and thus are removed.

```
In [26]: numeric_features = ["host_listings_count", "host_total_listings_count", "lat
percentage_features = ["host_response_rate", "host_acceptance_rate"]
currency_features = ["price"]
binary_features = ["host_is_superhost", "instant_bookable",]
categorical_features = ["name", "host_response_time", "neighbourhood_cleanse
# vector_features = ["neighborhood_overview",]

# drop_features = ["id", "host_id", "host_name", "neighbourhood"]
drop_features = ["id", "neighborhood_overview", "host_id", "host_name", "neig

numeric_transformer = make_pipeline(
    SimpleImputer(strategy="mean"),
    StandardScaler(),
)
percentage_transformer = make_pipeline(
    SimpleImputer(strategy="most_frequent"), # Since most data is string bet
    FunctionTransformer(strtopercent, validate=False),
    StandardScaler(),
```

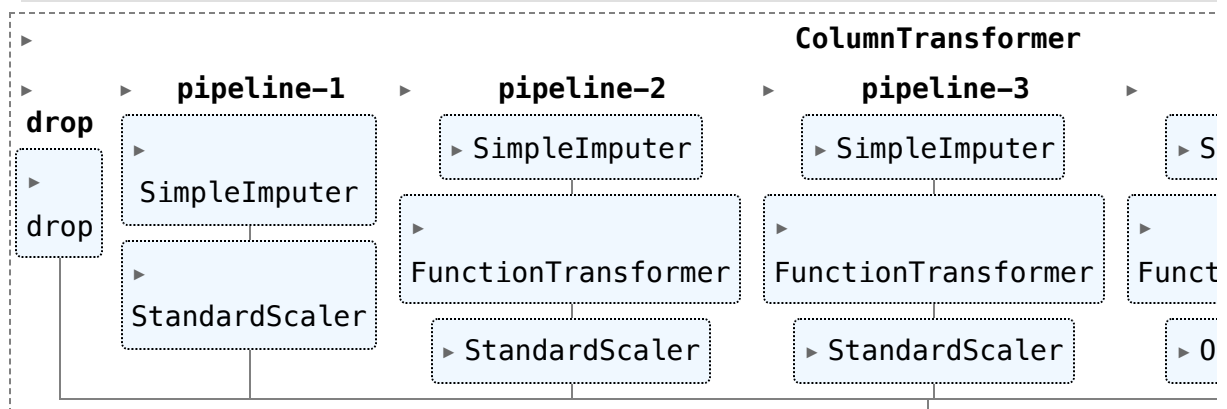


```

)
currency_transformer = make_pipeline(
    SimpleImputer(strategy="most_frequent"), # Since most data is string bet
    FunctionTransformer(strtomoney, validate=False),
    StandardScaler(),
)
binary_transformer = make_pipeline(
    SimpleImputer(strategy="most_frequent"),
    fun_tr,
    OneHotEncoder(handle_unknown="ignore", sparse_output=False, drop = "if_b
)
categorical_transformer = make_pipeline(
    SimpleImputer(strategy="most_frequent"),
    fun_tr,
    OneHotEncoder(handle_unknown="infrequent_if_exist", sparse_output=False)
)
# vector_transformer = make_pipeline(
#     CountVectorizer(stop_words="english"),
#     LogisticRegression(random_state=123)
# )

preprocessor = make_column_transformer(
    ("drop", drop_features),
    (numeric_transformer, numeric_features),
    (percentage_transformer, percentage_features),
    (currency_transformer, currency_features),
    (binary_transformer, binary_features),
    (categorical_transformer, categorical_features),
    # (vector_transformer, vector_features),
)
preprocessor.fit(X_train_transformed)
display(preprocessor)

```



```

In [27]: categorical_features_new = list(preprocessor.named_transformers_["pipeline-5
new_columns = (
    numeric_features + percentage_features + currency_features + binary_feat
)
X_train_enc = pd.DataFrame(
    preprocessor.transform(X_train_transformed), index=X_train_transformed.i
)
X_train_enc.head()

```

Out [27]:

	host_listings_count	host_total_listings_count	latitude	longitude	accommodates
0	-0.195679	-0.246259	0.996106	1.472784	-0.743460
1	-0.239566	-0.230072	-1.949916	0.522785	0.724263
2	-0.239566	-0.278633	-0.486575	0.367678	1.213504
3	-0.239566	-0.278633	-1.922130	0.378483	-0.743460
4	-0.217623	-0.246259	-0.163904	-1.319147	0.235022

5 rows × 143 columns

## Cross-validation Across Various Regression Models

- For optimisation of regression models, will use  $R^2$  as statistical measure of model quality to measure effectiveness in predicting future values.
- `mean_std_cross_val_scores()` provides both the mean and the standard deviation of the cross validations, allowing for a better understanding of the model's effectiveness and consistency.

```
In [28]: from sklearn.model_selection import (
    GridSearchCV,
    RandomizedSearchCV,
    cross_val_score,
    cross_validate,
    train_test_split,
)

def mean_std_cross_val_scores(model, X_train, y_train, **kwargs):
    """
    Returns mean and std of cross validation

    Parameters
    -----
    model :
        scikit-learn model
    X_train : numpy array or pandas DataFrame
        X in the training data
    y_train :
        y in the training data

    Returns
    -----
        pandas Series with mean scores from cross_validation
    """

    scores = cross_validate(model, X_train, y_train, **kwargs)

    mean_scores = pd.DataFrame(scores).mean()
```

```

std_scores = pd.DataFrame(scores).std()
out_col = []

for i in range(len(mean_scores)):
    out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i]))

return pd.Series(data=out_col, index=mean_scores.index)

```

```

In [29]: from sklearn.linear_model import Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, HistGradientBoostingRegressor

results = {}
scoring_metric = ["neg_root_mean_squared_error", "r2", "neg_mean_absolute_percentage_error"]

```

## Initial ridge, decision tree, random forest and tree-based ensemble models

```

In [30]: # Ridge regression model
# Hyperparameters: alpha
pipe_ridge = make_pipeline(pipeline_transformer, preprocessor, Ridge(alpha = 1e-5))
# pipe_ridge.fit(X_train, y_train)

results["Ridge"] = mean_std_cross_val_scores(
    pipe_ridge, X_train, y_train, cv = 10, return_train_score = True, scoring=scoring_metric
)

# Decision tree
# Hyperparameters: max_depth
# https://ken-hoffman.medium.com/decision-tree-hyperparameters-explained-491e1e1e1e1e
pipe_dt = make_pipeline(pipeline_transformer, preprocessor, DecisionTreeRegressor(max_depth=10))
# pipe_dt.fit(X_train, y_train)
results["Decision tree"] = mean_std_cross_val_scores(
    pipe_dt, X_train, y_train, cv = 10, return_train_score = True, scoring=scoring_metric
)

# Random forest
# Hyperparameters: n_estimators, max_depth, max_features
pipe_rf = make_pipeline(
    pipeline_transformer, preprocessor, RandomForestRegressor(n_jobs=-1, random_state=76))
# pipe_rf.fit(X_train, y_train)
results["Random forest"] = mean_std_cross_val_scores(
    pipe_rf, X_train, y_train, cv = 10, return_train_score = True, scoring=scoring_metric
)

# Tree-based ensemble model (HistGradientBoostingClassifier due to large dataset)
# Hyperparameters: learning_rate, max_depth
pipe_sklearn_histGB = make_pipeline(
    pipeline_transformer,
    preprocessor,
    HistGradientBoostingRegressor(random_state=76),
)
# pipe_sklearn_histGB.fit(X_train, y_train)

```

```

results["sklearn_histGB"] = mean_std_cross_val_scores(
    pipe_sklearn_histGB, X_train, y_train, cv = 10, return_train_score = True
)
# pd.DataFrame(results).T

```

```

/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))

```

## Ridge regression hyperparameter optimisation

```

In [31]: # Ridge regression model
# Hyperparameters: alpha
param_grid_ridge = {
    "ridge__alpha": [0.01, 0.1, 1, 10, 100],
}
random_search_ridge = RandomizedSearchCV(
    pipe_ridge, param_distributions = param_grid_ridge, n_iter=10,
    n_jobs = -1, return_train_score = True, random_state = 76, scoring='r2'
)
random_search_ridge.fit(X_train, y_train)

best_params_ridge = random_search_ridge.best_params_
best_score_ridge = random_search_ridge.best_score_

print("best_params_ridge: " + str(best_params_ridge))
print("best_score_ridge: " + str(best_score_ridge))

pipe_ridge_opt = make_pipeline(
    pipeline_transformer,
    preprocessor,
    Ridge(
        alpha = best_params_ridge["ridge__alpha"],
        max_iter = 200,
    )
)
pipe_ridge_opt.fit(X_train, y_train)

```

```
results["Ridge (Optimised)"] = mean_std_cross_val_scores(pipe_ridge_opt, X_t
# pd.DataFrame(results).T
```

```
/opt/conda/lib/python3.10/site-packages/sklearn/model_selection/_search.py:3
05: UserWarning: The total space of parameters 5 is smaller than n_iter=10.
Running 5 iterations. For exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
best_params_ridge: {'ridge__alpha': 100}
best_score_ridge: 0.2835346371356892
```

```
/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treati
ng keys as positions is deprecated. In a future version, integer keys will a
lways be treated as labels (consistent with DataFrame behavior). To access a
value by position, use `ser.iloc[pos]`
out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
```

## Decision tree hyperparameter optimisation

```
In [32]: # Decision tree
# Hyperparameters: max_depth
# https://ken-hoffman.medium.com/decision-tree-hyperparameters-explained-491
param_grid_dt = {
    "decisiontreeregressor__max_depth": np.arange(1, 20, 4),
}
random_search_dt = RandomizedSearchCV(
    pipe_dt, param_distributions = param_grid_dt, n_iter=10,
    n_jobs = -1, return_train_score = True, random_state = 76, scoring='r2'
)
random_search_dt.fit(X_train, y_train)

best_params_dt = random_search_dt.best_params_
best_score_dt = random_search_dt.best_score_
print("best_params_dt: " + str(best_params_dt))
print("best_score_dt: " + str(best_score_dt))

pipe_dt_opt = make_pipeline(
    pipeline_transformer,
    preprocessor,
    DecisionTreeRegressor(
        max_depth = best_params_dt["decisiontreeregressor__max_depth"],
        random_state=76
    )
)
results["Decision tree (Optimised)"] = mean_std_cross_val_scores(
    pipe_dt_opt, X_train, y_train, cv = 10, return_train_score = True, scorin
)
```

```
/opt/conda/lib/python3.10/site-packages/sklearn/model_selection/_search.py:3
05: UserWarning: The total space of parameters 5 is smaller than n_iter=10.
Running 5 iterations. For exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
best_params_dt: {'decisiontreeregressor__max_depth': 5}
best_score_dt: 0.242331494810347
```

```
/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
```

## Random forest hyperparameter optimisation

```
In [33]: # Random forest
# Hyperparameters: n_estimators, max_depth, max_features, class_weight
param_grid_rf = {
    "randomforestregressor__max_depth": np.arange(1, 20, 4),
    "randomforestregressor__n_estimators": [100, 200, 300, 400, 500],
    "randomforestregressor__max_features": [10, 15, 20, 25, 30, 35, 40],
}
random_search_rf = RandomizedSearchCV(
    pipe_rf, param_distributions = param_grid_rf, n_iter=10,
    n_jobs = -1, return_train_score = True, random_state = 76, scoring='r2'
)
random_search_rf.fit(X_train, y_train)

best_params_rf = random_search_rf.best_params_
best_score_rf = random_search_rf.best_score_
print("best_params_rf: " + str(best_params_rf))
print("best_score_rf: " + str(best_score_rf))

pipe_rf_opt = make_pipeline(
    pipeline_transformer,
    preprocessor,
    RandomForestRegressor(
        n_estimators = best_params_rf["randomforestregressor__n_estimators"],
        max_depth = best_params_rf["randomforestregressor__max_depth"],
        max_features = best_params_rf["randomforestregressor__max_features"],
        n_jobs=-1, random_state=76,
    )
)
results["Random forest (Optimised)"] = mean_std_cross_val_scores(
    pipe_rf_opt, X_train, y_train, cv = 10, return_train_score = True, scoring='r2'
)
```

```
best_params_rf: {'randomforestregressor__n_estimators': 400, 'randomforestregressor__max_features': 35, 'randomforestregressor__max_depth': 5}
best_score_rf: 0.2854408171022718
```

```
/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
  out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))
```

## Tree-based ensemble hyperparameter optimisation

```
In [34]: # Tree-based ensemble model (HistGradientBoostingRegressor due to large data)
# Hyperparameters: learning_rate, max_depth
param_grid_sklearn_histGB = {
```

```

        "histgradientboostingregressor__max_depth": np.arange(1, 20, 4),
        "histgradientboostingregressor__learning_rate": [0.1, 0.3, 0.5, 0.7, 0.9]
    }
    random_search_sklearn_histGB = RandomizedSearchCV(
        pipe_sklearn_histGB, param_distributions = param_grid_sklearn_histGB, n_
        n_jobs = -1, return_train_score = True, random_state = 76, scoring='r2'
    )
    random_search_sklearn_histGB.fit(X_train, y_train)

    best_params_sklearn_histGB = random_search_sklearn_histGB.best_params_
    best_score_sklearn_histGB = random_search_sklearn_histGB.best_score_

    print("best_params_sklearn_histGB: " + str(best_params_sklearn_histGB))
    print("best_score_sklearn_histGB: " + str(best_score_sklearn_histGB))

    pipe_sklearn_histGB_opt = make_pipeline(
        pipeline_transformer,
        preprocessor,
        HistGradientBoostingRegressor(
            learning_rate = best_params_sklearn_histGB["histgradientboostingregressor__learning_rate"],
            max_depth = best_params_sklearn_histGB["histgradientboostingregressor__max_depth"],
            random_state=76
        ),
    )
    results["sklearn_histGB (Optimised)"] = mean_std_cross_val_scores(
        pipe_sklearn_histGB_opt, X_train, y_train, cv = 10, return_train_score = True
    )

```

```

best_params_sklearn_histGB: {'histgradientboostingregressor__max_depth': 1,
'histgradientboostingregressor__learning_rate': 0.3}
best_score_sklearn_histGB: 0.29169275617122853

```

```

/tmp/ipykernel_30/3623116645.py:34: FutureWarning: Series.__getitem__ treating
keys as positions is deprecated. In a future version, integer keys will al
ways be treated as labels (consistent with DataFrame behavior). To access a
value by position, use `ser.iloc[pos]`
    out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores[i], std_scores[i])))

```

## Results from All Models

- `pipe_dt` is likely to have overfitted on the training data as its `train_r2` score is unrealistically high, and its `test_r2` score dropped into the negatives.
- `pipe_sklearn_histGB_opt` provides the greatest `test_r2`, while `pipe_rf_opt` provides a similar `test_r2` with a smaller standard deviation.
- `test_neg_root_mean_squared_error` and `test_neg_mean_absolute_percentage_error` have smaller absolute values for `pipe_sklearn_histGB_opt`.
- As such, `pipe_sklearn_histGB_opt` is arguably slightly more optimal for prediction, along with requiring less time for fitting and scoring.

```
In [35]: pd.DataFrame(results).T.sort_values('test_r2', ascending=False)
```

Out[35]:

	fit_time	score_time	test_neg_root_mean_squared_error	train_neg_rc
<b>sklearn_histGB (Optimised)</b>	0.255 (+/- 0.006)	0.031 (+/- 0.001)	-1120.762 (+/- 33.061)	
<b>Random forest (Optimised)</b>	1.433 (+/- 0.105)	0.126 (+/- 0.006)	-1125.229 (+/- 35.699)	
<b>Ridge (Optimised)</b>	0.243 (+/- 0.065)	0.052 (+/- 0.016)	-1127.742 (+/- 34.152)	
<b>Ridge</b>	0.248 (+/- 0.073)	0.047 (+/- 0.005)	-1131.805 (+/- 34.525)	
<b>Random forest</b>	3.206 (+/- 0.104)	0.060 (+/- 0.001)	-1153.308 (+/- 28.898)	
<b>sklearn_histGB</b>	0.955 (+/- 0.029)	0.032 (+/- 0.001)	-1155.255 (+/- 30.932)	
<b>Decision tree (Optimised)</b>	0.138 (+/- 0.002)	0.024 (+/- 0.001)	-1168.827 (+/- 44.247)	
<b>Decision tree</b>	0.237 (+/- 0.027)	0.024 (+/- 0.001)	-1629.723 (+/- 33.542)	

## Production of Results

```
In [36]: # Given test_r2, pipe_rf_opt provides the best prediction
pipe_sklearn_histGB_opt.fit(X_train, y_train)
y_preds = pd.DataFrame(pipe_sklearn_histGB_opt.predict(test_df).T, columns=
ids = pd.DataFrame(test_df["id"], columns=['id'])

df = pd.concat([ids, y_preds], axis=1)
# df.head()
df.info()
df.to_csv('submission.csv', index=False)
print('Generated Submission file')
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1339 entries, 0 to 1338
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               1339 non-null   int64
1   monthly_revenue  1339 non-null   float64
dtypes: float64(1), int64(1)
memory usage: 21.0 KB
Generated Submission file
```